

Kent Gener

Kent Hervey D. Gener BSIT2-B

2023-11-22

R Markdown

This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com>.

When you click the **Knit** button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. You can embed an R code chunk like this:

```
summary(cars)
```

```
##      speed      dist
##  Min.   : 4.0    Min.   :  2.00
##  1st Qu.:12.0    1st Qu.: 26.00
##  Median :15.0    Median : 36.00
##  Mean   :15.4    Mean   : 42.98
##  3rd Qu.:19.0    3rd Qu.: 56.00
##  Max.   :25.0    Max.   :120.00
```

Including Plots

You can also embed plots, for example:



Note that the `echo = FALSE` parameter was added to the code chunk to prevent printing of the R code that generated the plot.

- Download and open the mpg file. Upload it to your OWN environment a. Show your solutions on how to import a csv file into the environment.

```
mpg_data <- read.csv('mpg.csv')
```

- b. Which variables from mpg dataset are categorical?

```
str(mpg_data)
```

```
## 'data.frame':  234 obs. of  12 variables:
## $ X          : int  1 2 3 4 5 6 7 8 9 10 ...
## $ manufacturer: chr  "audi" "audi" "audi" "audi" ...
## $ model       : chr  "a4" "a4" "a4" "a4" ...
## $ displ       : num  1.8 1.8 2 2 2.8 2.8 3.1 1.8 1.8 2 ...
## $ year        : int  1999 1999 2008 2008 1999 1999 2008 1999 1999 2008 ...
## $ cyl         : int  4 4 4 4 6 6 6 4 4 4 ...
## $ trans       : chr  "auto(l5)" "manual(m5)" "manual(m6)" "auto(av)" ...
## $ drv         : chr  "f" "f" "f" "f" ...
## $ cty         : int  18 21 20 21 16 18 18 18 16 20 ...
## $ hwy         : int  29 29 31 30 26 26 27 26 25 28 ...
## $ fl         : chr  "p" "p" "p" "p" ...
## $ class       : chr  "compact" "compact" "compact" "compact" ...
```

```
categorical_vars <- c("manufacturer", "model", "year", "cyl", "trans", "drv", "fl", "class")
cat("Categorical variables:", categorical_vars, "\n")
```

```
## Categorical variables: manufacturer model year cyl trans drv fl class
```

c. Which are continuous variables?

```
str(mpg_data)
```

```
## 'data.frame': 234 obs. of 12 variables:
## $ X : int 1 2 3 4 5 6 7 8 9 10 ...
## $ manufacturer: chr "audi" "audi" "audi" "audi" ...
## $ model : chr "a4" "a4" "a4" "a4" ...
## $ displ : num 1.8 1.8 2 2 2.8 2.8 3.1 1.8 1.8 2 ...
## $ year : int 1999 1999 2008 2008 1999 1999 2008 1999 1999 2008 ...
## $ cyl : int 4 4 4 4 6 6 6 4 4 4 ...
## $ trans : chr "auto(l5)" "manual(m5)" "manual(m6)" "auto(av)" ...
## $ drv : chr "f" "f" "f" "f" ...
## $ cty : int 18 21 20 21 16 18 18 16 20 ...
## $ hwy : int 29 29 31 30 26 26 27 26 25 28 ...
## $ fl : chr "p" "p" "p" "p" ...
## $ class : chr "compact" "compact" "compact" "compact" ...
```

```
continuous_vars <- c("displ", "cty", "hwy")
cat("Continuous variables:", continuous_vars, "\n")
```

```
## Continuous variables: displ cty hwy
```

2. Which manufacturer has the most models in this data set? Which model has the most variations?
Show your answer.

```
library(dplyr)
```

```
## Warning: package 'dplyr' was built under R version 4.3.2
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
## filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
## intersect, setdiff, setequal, union
```

```
models_per_manufacturer <- mpg_data %>%
  group_by(manufacturer) %>%
  summarise(unique_models = n_distinct(model))
```

```
models_per_manufacturer
```

```
## # A tibble: 15 x 2
##   manufacturer unique_models
##   <chr>          <int>
## 1 audi           3
## 2 chevrolet      4
## 3 dodge          4
## 4 ford           4
## 5 honda          1
## 6 hyundai        2
## 7 jeep           1
## 8 land rover     1
## 9 lincoln        1
## 10 mercury        1
## 11 nissan          3
## 12 pontiac        1
## 13 subaru         2
## 14 toyota         6
## 15 volkswagen     4
```

a. Group the manufacturers and find the unique models. Show your codes and result.

```
library(dplyr)

models_per_manufacturer <- mpg_data %>%
  group_by(manufacturer) %>%
  summarise(unique_models = n_distinct(model))

models_per_manufacturer
```

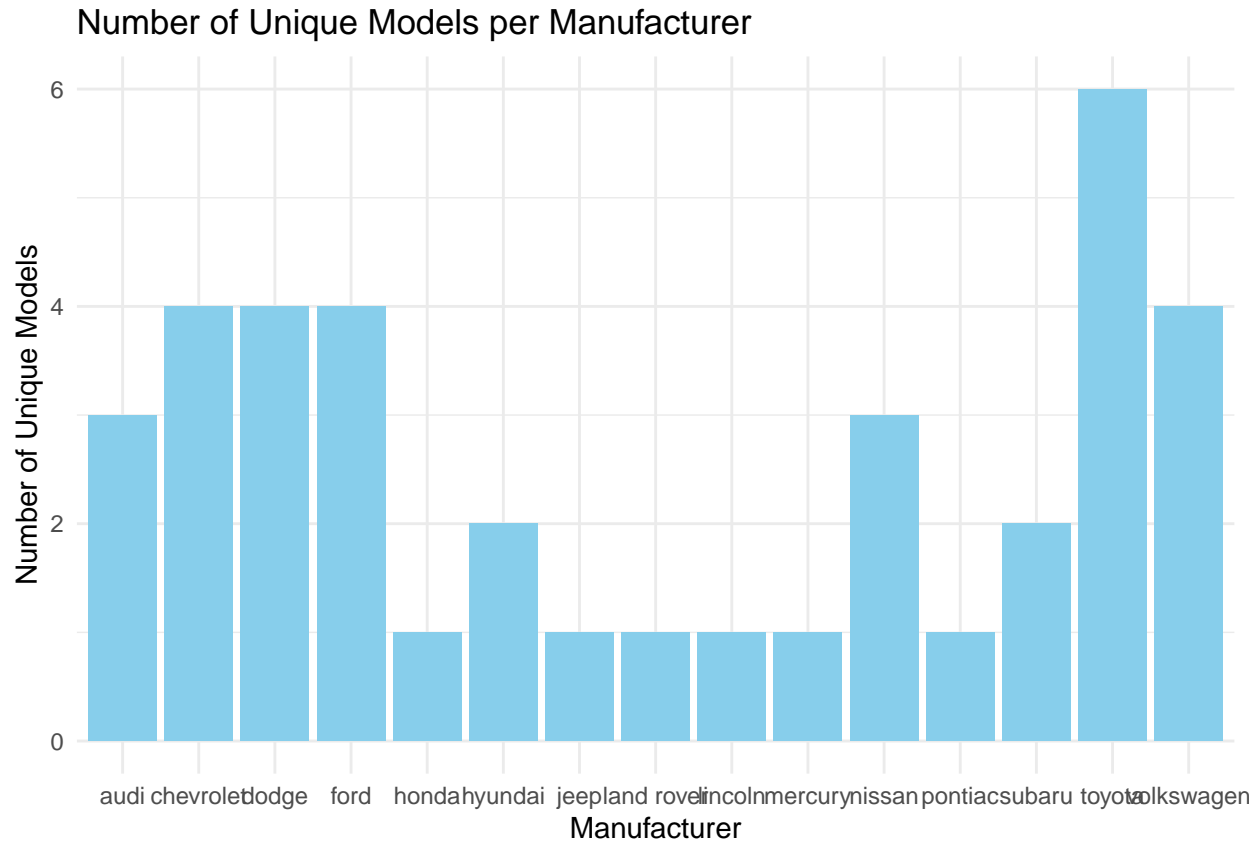
```
## # A tibble: 15 x 2
##   manufacturer unique_models
##   <chr>          <int>
## 1 audi           3
## 2 chevrolet      4
## 3 dodge          4
## 4 ford           4
## 5 honda          1
## 6 hyundai        2
## 7 jeep           1
## 8 land rover     1
## 9 lincoln        1
## 10 mercury        1
## 11 nissan          3
## 12 pontiac        1
## 13 subaru         2
## 14 toyota         6
## 15 volkswagen     4
```

b. Graph the result by using plot() and ggplot(). Write the codes and its result.

```
library(ggplot2)
```

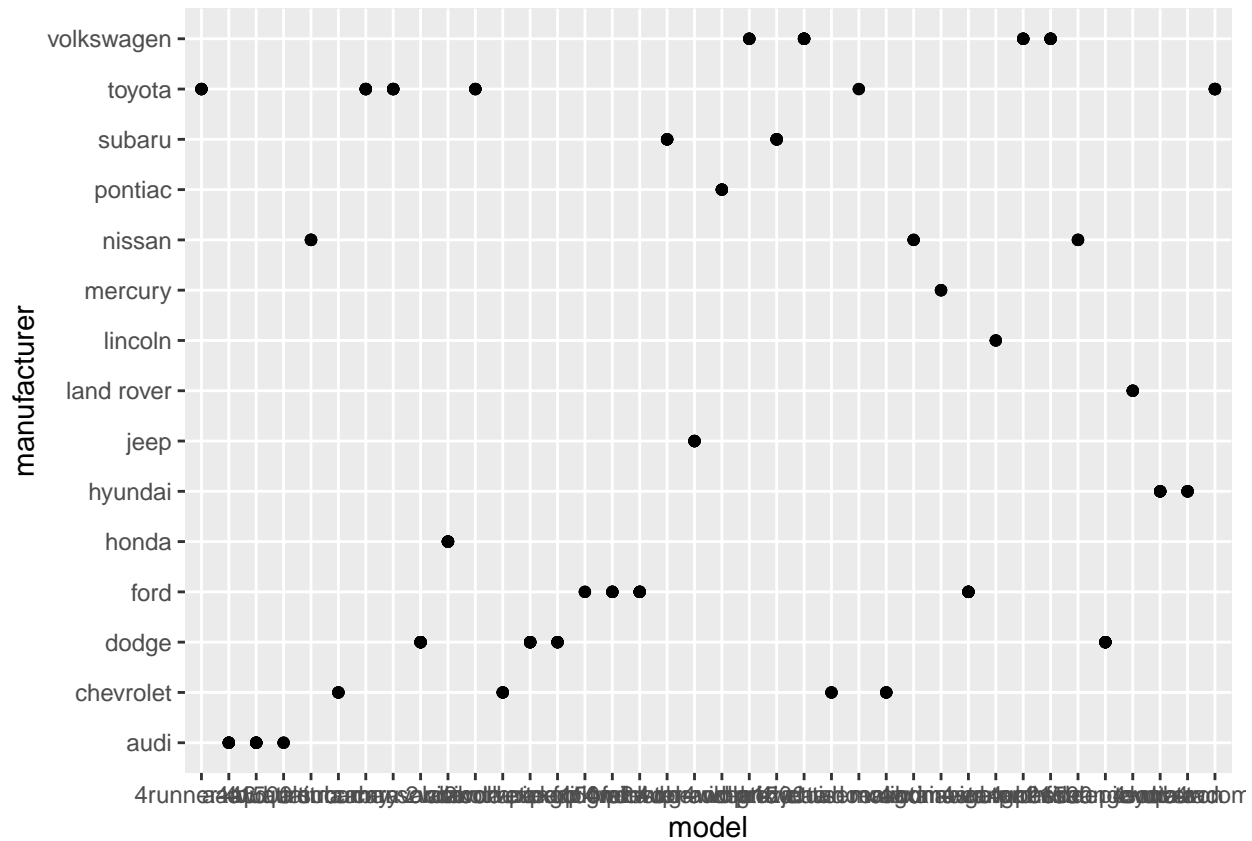
```
## Warning: package 'ggplot2' was built under R version 4.3.2
```

```
ggplot(models_per_manufacturer, aes(x = manufacturer, y = unique_models)) +
  geom_bar(stat = "identity", fill = "skyblue") +
  labs(title = "Number of Unique Models per Manufacturer", x = "Manufacturer", y = "Number of Unique Models") +
  theme_minimal()
```



2. Same dataset will be used. You are going to show the relationship of the model and the manufacturer.
 - a. What does `ggplot(mpg, aes(model, manufacturer)) + geom_point()` show?

```
ggplot(mpg_data, aes(model, manufacturer)) + geom_point()
```




b. For you, is it useful? If not, how could you modify the data to make it more informative?

```
library(ggplot2)
```

```
model_manufacturer_count <- mpg %>%
  group_by(model, manufacturer) %>%
  summarise(count = n())
```

```
## 'summarise()' has grouped output by 'model'. You can override using the
## '.groups' argument.
```

```
ggplot(model_manufacturer_count, aes(x = model, fill = manufacturer)) +
  geom_bar() +
  labs(title = "Number of Models per Manufacturer",
       x = "Model", y = "Count",
       fill = "Manufacturer") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

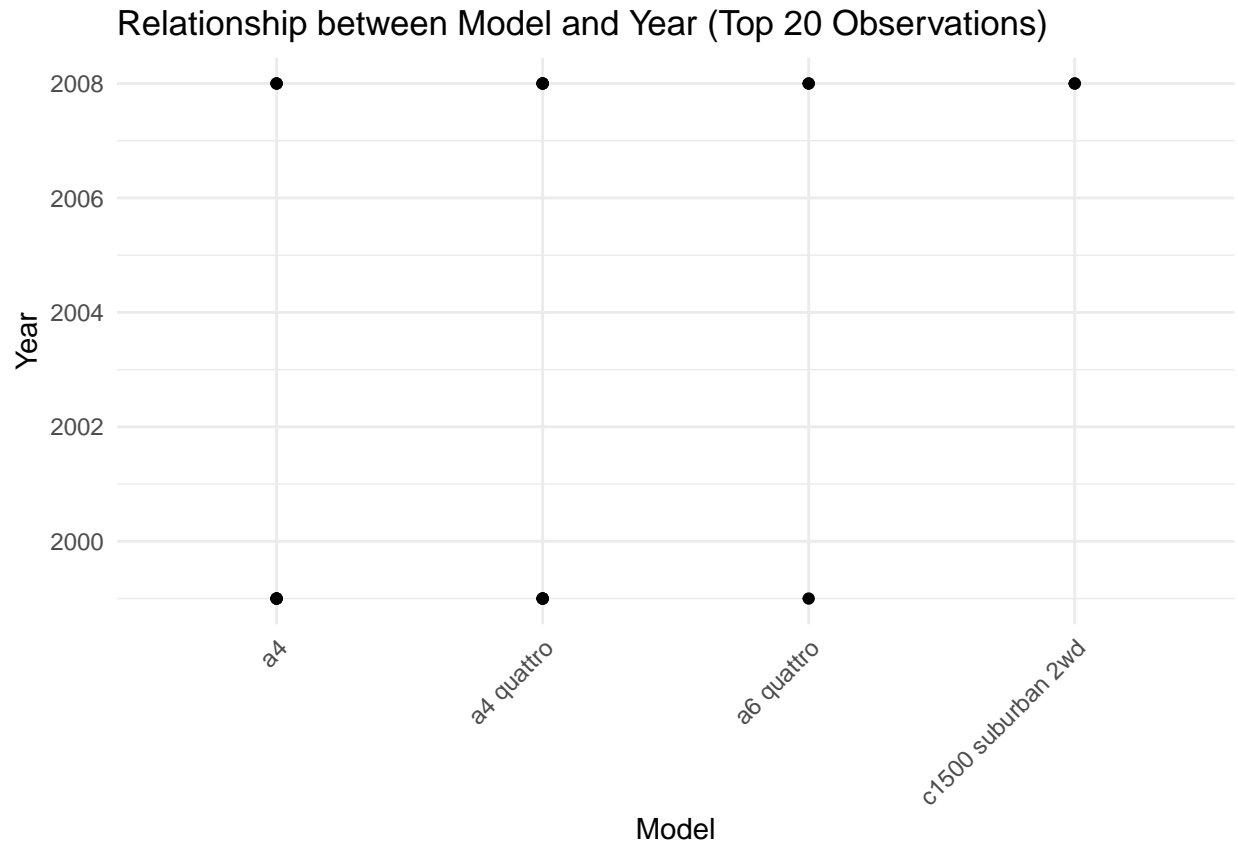


3. Plot the model and the year using ggplot(). Use only the top 20 observations. Write the codes and its results.

```
library(ggplot2)

top_20 <- head(mpg, 20)

ggplot(top_20, aes(x = model, y = year)) +
  geom_point() +
  labs(title = "Relationship between Model and Year (Top 20 Observations)",
       x = "Model", y = "Year") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```



4. Using the pipe (`%>%`), group the model and get the number of cars per model. Show codes and its result

```
library(dplyr)

cars_per_model <- mpg %>%
  group_by(model) %>%
  summarise(num_cars = n())

cars_per_model
```

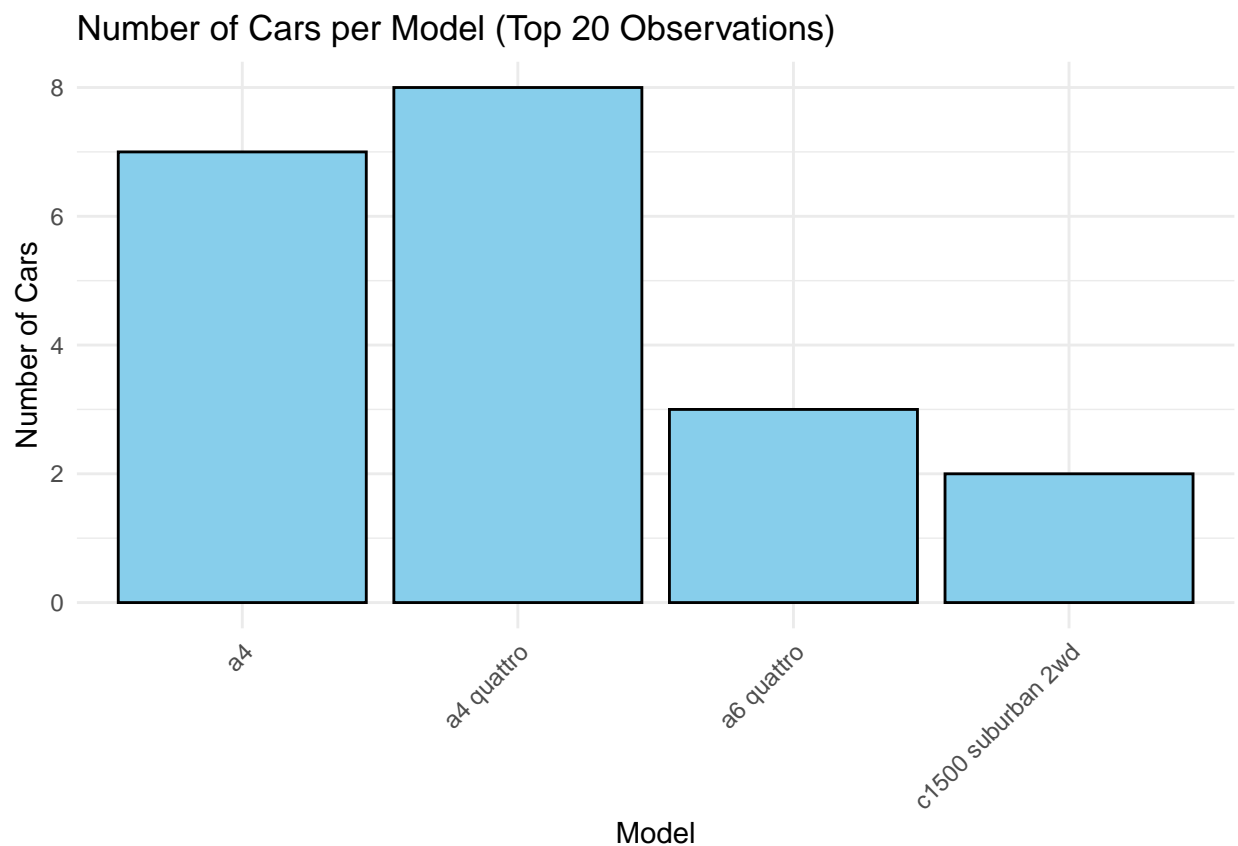
```
## # A tibble: 38 x 2
##   model          num_cars
##   <chr>          <int>
## 1 4runner 4wd             6
## 2 a4                     7
## 3 a4 quattro             8
## 4 a6 quattro             3
## 5 altima                 6
## 6 c1500 suburban 2wd     5
## 7 camry                  7
## 8 camry solara           7
## 9 caravan 2wd           11
## 10 civic                  9
## # i 28 more rows
```


- a. Plot using `geom_bar()` using the top 20 observations only. The graphs should have a title, labels and colors. Show code and results.

```
library(ggplot2)

top_20 <- head(mpg, 20)

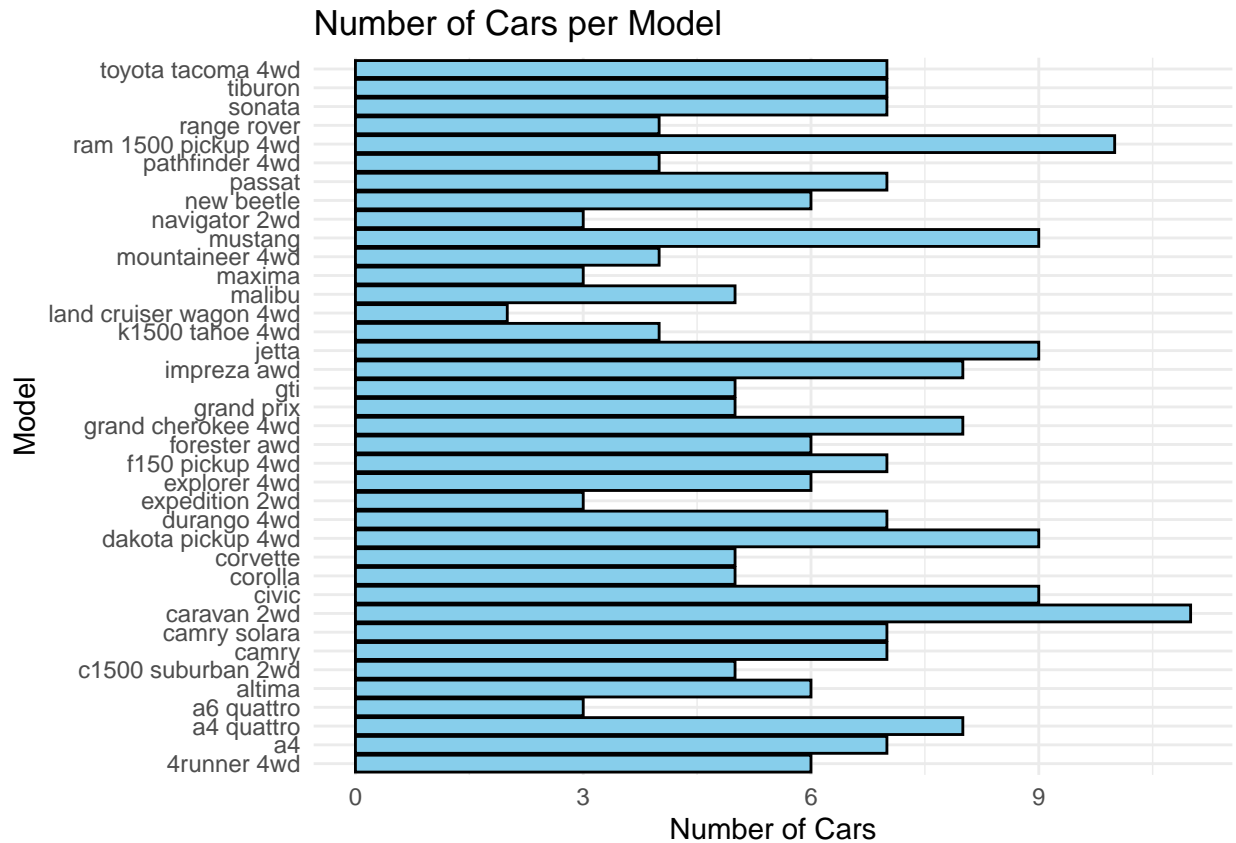
ggplot(top_20, aes(x = model)) +
  geom_bar(fill = "skyblue", color = "black") +
  labs(title = "Number of Cars per Model (Top 20 Observations)",
       x = "Model", y = "Number of Cars") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```



- b. Plot using the `geom_bar()` + `coord_flip()` just like what is shown below. Show codes and its result.

```
library(ggplot2)


ggplot(cars_per_model, aes(x = model, y = num_cars)) +
  geom_bar(fill = "skyblue", color = "black", stat = "identity") +
  labs(title = "Number of Cars per Model",
       x = "Model", y = "Number of Cars") +
  coord_flip() +
  theme_minimal()
```



5. Plot the relationship between `cyl` - number of cylinders and `displ` - engine displacement using `geom_point` with aesthetic `color = engine displacement`. Title should be "Relationship between No. of Cylinders and Engine Displacement".

```
library(ggplot2)

ggplot(mpg, aes(x = cyl, y = displ, color = displ)) +
  geom_point() +
  labs(title = "Relationship between No. of Cylinders and Engine Displacement",
       x = "Number of Cylinders", y = "Engine Displacement") +
  scale_color_continuous(name = "Engine Displacement") +
  theme_minimal()
```



a. How would you describe its relationship? Show the codes and its result.

#The relationship between cylinder count and engine displacement is often positive, with larger engine

6. Plot the relationship between displ (engine displacement) and hwy(highway miles per gallon). Mapped it with a continuous variable you have identified in #1-c. What is its result? Why it produced such output?

7. Import the traffic.csv onto your R environment.

```
traffic <- read.csv("traffic.csv")
head(traffic)
```

```
##           DateTime Junction Vehicles          ID
## 1 2015-11-01 00:00:00         1      15 20151101001
## 2 2015-11-01 01:00:00         1      13 20151101011
## 3 2015-11-01 02:00:00         1      10 20151101021
## 4 2015-11-01 03:00:00         1       7 20151101031
## 5 2015-11-01 04:00:00         1       9 20151101041
## 6 2015-11-01 05:00:00         1       6 20151101051
```

a. How many numbers of observation does it have? What are the variables of the traffic dataset the Show your answer.

```

num_observations <- nrow(traffic)

variables <- names(traffic)

cat("Number of Observations:", num_observations, "\n")

## Number of Observations: 48120

cat("Variables:", paste(variables, collapse = ", "), "\n")

## Variables: DateTime, Junction, Vehicles, ID

```

b. subset the traffic dataset into junctions. What is the R codes and its output?

```

junctions <- unique(traffic$junction)

junctions

## NULL

```

c. Plot each junction in a using geom_line(). Show your solution and output.

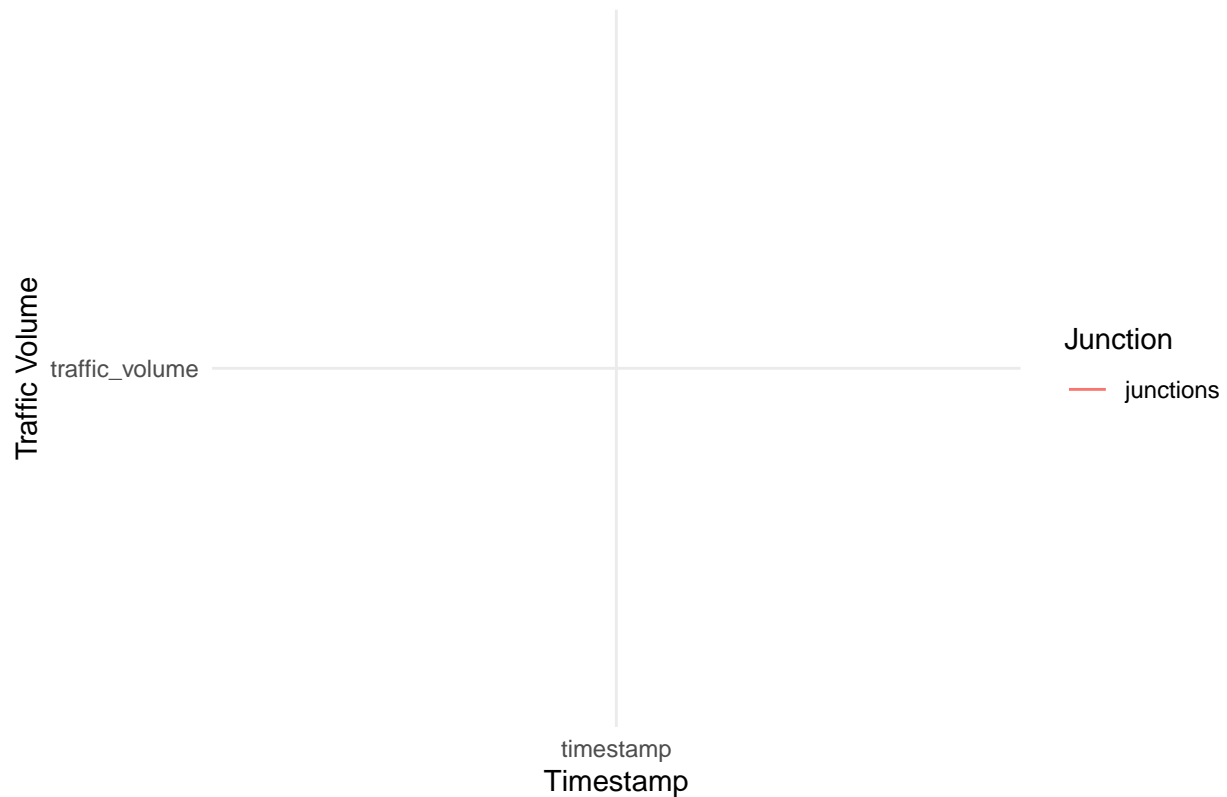
```

library(ggplot2)

# Plot each junction using geom_line()
ggplot(traffic, aes(x = 'timestamp', y = 'traffic_volume', color = 'junctions')) +
  geom_line() +
  labs(title = "Traffic Volume Over Time for Each Junction",
       x = "Timestamp", y = "Traffic Volume",
       color = "Junction") +
  theme_minimal()

```

Traffic Volume Over Time for Each Junction



7. From alexa_file.xlsx, import it to your environment

```
if (!requireNamespace("readxl", quietly = TRUE)) {
  install.packages("readxl")
}
library(readxl)
```

```
## Warning: package 'readxl' was built under R version 4.3.2
```

```
file_path <- "alexa_file.xlsx"

alexa_data <- read_excel(file_path)

head(alexa_data)
```

```
## # A tibble: 6 x 5
##   rating date          variation      verified_reviews      feedback
##   <dbl> <dtm>          <chr>          <chr>          <dbl>
## 1     5 2018-07-31 00:00:00 Charcoal Fabric Love my Echo!          1
## 2     5 2018-07-31 00:00:00 Charcoal Fabric Loved it!              1
## 3     4 2018-07-31 00:00:00 Walnut Finish  Sometimes while playi~ 1
## 4     5 2018-07-31 00:00:00 Charcoal Fabric I have had a lot of f~ 1
## 5     5 2018-07-31 00:00:00 Charcoal Fabric Music                1
## 6     5 2018-07-31 00:00:00 Heather Gray Fabric I received the echo a~ 1
```

- a. How many observations does alexa_file has? What about the number of columns? Show your solution and answer.

```
num_observations <- nrow(alexa_data)

num_columns <- ncol(alexa_data)

cat("Number of Observations:", num_observations, "\n")
```

```
## Number of Observations: 3150
```

```
cat("Number of Columns:", num_columns, "\n")
```

```
## Number of Columns: 5
```

- b. group the variations and get the total of each variations. Use dplyr package. Show solution and answer.

```
library(dplyr)

variation_totals <- alexa_data %>%
  group_by(variation) %>%
  summarise(total = n())

variation_totals
```

```
## # A tibble: 16 x 2
##   variation          total
##   <chr>          <int>
## 1 Black          261
## 2 Black Dot      516
## 3 Black Plus     270
## 4 Black Show     265
## 5 Black Spot     241
## 6 Charcoal Fabric 430
## 7 Configuration: Fire TV Stick 350
## 8 Heather Gray Fabric 157
## 9 Oak Finish      14
## 10 Sandstone Fabric 90
## 11 Walnut Finish   9
## 12 White          91
## 13 White Dot      184
## 14 White Plus      78
## 15 White Show     85
## 16 White Spot     109
```

- c. Plot the variations using the ggplot() function. What did you observe? Complete the details of the graph. Show solution and answer.

```
library(ggplot2)

ggplot(alexa_data, aes(x = variation)) +
```

```
geom_bar() +
labs(title = "Distribution of Alexa Variations",
     x = "Variation",
     y = "Count") +
theme_minimal()
```

